

HW04

May 16, 2020

1 Computer Vision - HW04 - 98722278

Index:

1. Canny Edge Detector
 1. Gaussian Noise
 2. Gradient Intensification
 3. Non-Max Suppression
 4. Thresholding
2. Straight Line Detector
 1. Assign each edge to a direction
 2. Getting edgelets using connected components
 3. Compute straightness and theta
 4. Threshold
 5. Test

1.1 1 Implementation of Canny Edge Detector

1. Gaussian Noise
2. Gradient Intensification
3. Non-Max Suppression
4. Thresholding

```
In [1]: import cv2
import numpy as np
from scipy import ndimage
from PIL import Image

import matplotlib.pyplot as plt

import time
%matplotlib inline

In [2]: def open_image(path):
        """
        Open an image using PIL library
```

```

        :param path: path to image file-like
        :return: PIL image object
        """
        image = Image.open(path)
        return image

In [3]: def show_image(image, cmap='gray'):
        """
        Show PIL image or numpy image in default viewer of OS

        :param image: image data
        :param cmap: color map of input numpy array
        :return: None
        """
        if str(type(image)).__contains__('PIL'):
            image.show()
        elif str(type(image)).__contains__('numpy'):
            if cmap=='gray':
                Image.fromarray(np.uint8(image), mode='L').show()
            elif cmap == 'bw':
                size = image.shape[::-1]
                data_bytes = np.packbits(image, axis=1)
                Image.frombytes(mode='1', size=size, data=data_bytes).show()
            else:
                raise ValueError('color map is invalid.')
        else:
            raise ValueError('Input t is not valid.')

In [4]: class ToGrayscale:
        """
        Get and PIL image or numpy n-dim array as image and convert it to grayscale image
        """

        def __init__(self):
            pass

        def __call__(self, image):
            """
            Get and PIL image or numpy n-dim array as image and convert it to grayscale image

            :param image: input image data
            :return: Grayscale image of input type
            """
            if str(type(image)).__contains__('PIL'):
                image = image.convert('L')
            elif str(type(image)).__contains__('numpy'):
                image = np.dot(image[..., :3], [0.2989, 0.5870, 0.1140])

```

```

else:
    raise ValueError('Input type is not valid.')
return image

```

1.1.1 1.A Gaussian Noise

```

In [5]: class GaussianNoise:
    def __init__(self, size=5, std=1):
        self.size = size
        self.std = std

    def _gaussian(self, r2):
        """
        Sample one instance from gaussian distribution regarding
        given squared-distance:r2, standard-deviation:std and general-constant:k

        :param r: squared distance from center of gaussian distribution
        :param std: standard deviation

        :return: A sampled number obtained from gaussian
        """
        return np.exp(-r2/(2.*self.std**2)) / (2.*np.pi*self.std**2)

    def _gaussian_kernel(self):
        """
        Creates a gaussian kernel regarding given size and std.
        Note that to define interval with respect to the size,
        I used linear space sampling which may has
        lower accuracy from renowned libraries.

        :param std: standard deviation value
        :param size: size of the output kernel
        :return: A gaussian kernel with size of (size*size)
        """
        self.size = int(self.size) // 2
        x, y = np.mgrid[-self.size:self.size+1, -self.size:self.size+1]
        distance = x**2+ y**2
        kernel = self._gaussian(r2=distance)
        return kernel

    def __call__(self, image):
        """
        Applies gaussian noise on the given image

        :param image: Input image in grayscale mode numpy ndarray or cv2 image
        :param size: Size of the gaussian kernel
        :param std: Standard deviation value for gaussian kernel
        """

```

```
return ndimage.convolve(image, self._gaussian_kernel())
```

```
In [9]: image = cv2.imread('images/cameraman.jpg', cv2.IMREAD_GRAYSCALE)
        gaussian_noise = GaussianNoise()
        image_blurred = gaussian_noise(image)
```

```
# plotting
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20, 15))
ax[0].set_title('original')
ax[1].set_title('blurred')
ax[0].imshow(image, cmap='gray')
ax[1].imshow(image_blurred, cmap='gray')
```

```
Out [9]: <matplotlib.image.AxesImage at 0x17c7aad22e8>
```



1.1.2 1.B Gradient Intensity

```
In [6]: class GradientIntensity:
```

```
    """
```

We use Sobel filters to convolve over image (numpy ndarray) to calculate gradient in horizontal and vertical directions. Finally returns magnitude G and slope θ as

$G = \sqrt{I_x^2 + I_y^2}$

$\theta = \arctan(I_x/I_y)$

We use these Sobel filters as default:

$K_x =$

```

[[ -1  0  1],
 [-2  0  2],
 [-1  0  1]]

Ky =
[[ 1  2  1],
 [ 0  0  0],
 [-1 -2 -1]]

"""

def __init__(self, hf=None, vf=None, init=True):
    """
    Initialize filters

    :param hf: Horizontal filter matrix -> numpy ndarray
    :param vf: Vertical filter matrix -> numpy ndarray
    :param init: whether initialize Sobel filters or initialize using user provided
    """

    if not init:
        self.hf = hf
        self.vf = vf
    else:
        self.hf = np.array(
            [[-1, 0, 1],
             [-2, 0, 2],
             [-1, 0, 1]])

        self.vf = np.array(
            [[1, 2, 1],
             [0, 0, 0],
             [-1, -2, -1]])

    def __call__(self, x):
        if not str(type(x)).__contains__('numpy'):
            raise ValueError('Invalid input. Please provide numpy ndarray image.')
        Ix = ndimage.filters.convolve(x, self.hf)
        Iy = ndimage.filters.convolve(x, self.vf)

        G = np.sqrt(np.power(Ix, 2) + np.power(Iy, 2))
        G = G / G.max() * 255
        theta = np.arctan2(Iy, Ix)

        return G, theta

```

```

In [11]: to_grayscale = ToGrayscale()
         image = np.array(to_grayscale(open_image('images/cameraman.jpg')), dtype=float) # th

```


```

gaussian_noise = GaussianNoise()
image_blurred = gaussian_noise(image)
gradient_intensity = GradientIntensity()
image_grad, image_theta = gradient_intensity(image_blurred)

# plotting
fig, ax = plt.subplots(nrows=1, ncols=3, figsize=(20, 15))
ax[0].set_title('blurred')
ax[1].set_title('grad')
ax[2].set_title('theta')
ax[0].imshow(image_blurred, cmap='gray')
ax[1].imshow(image_grad, cmap='gray')
ax[2].imshow(image_theta, cmap='gray')

```

Out[11]: <matplotlib.image.AxesImage at 0x17c7ac28550>



1.1.3 1.C Non-Max Suppression

In [7]: `class NonMaxSuppression:`

"""

Get gradient of image w.r.t the filters and degree of gradients (theta) and keep most intensified pixel in each direction.

Note: d_prime = d-180

"""

```

def __init__(self):
    pass

```

```

def __call__(self, grad_img, grad_dir):
    """

```

Get non-max suppressed image by preserving most intensified pixels

:param grad_img: Gradient image gathered by convolving filters on original image

:param grad_dir: Gradient directions gathered by convolving filters on original image

```

        :return: Soft-edge numpy ndarray image
        """

z = np.zeros(shape=grad_img.shape, dtype=np.int32)

for h in range(grad_img.shape[0]):
    for v in range(grad_img.shape[1]):
        degree = self.__angle__(grad_dir[h][v])

        try:
            if degree == 0:
                if grad_img[h][v] >= grad_img[h][v - 1] and grad_img[h][v] >= grad_img[h][v + 1]:
                    z[h][v] = grad_img[h][v]
            elif degree == 45:
                if grad_img[h][v] >= grad_img[h - 1][v + 1] and grad_img[h][v] >= grad_img[h + 1][v - 1]:
                    z[h][v] = grad_img[h][v]
            elif degree == 90:
                if grad_img[h][v] >= grad_img[h - 1][v] and grad_img[h][v] >= grad_img[h + 1][v]:
                    z[h][v] = grad_img[h][v]
            elif degree == 135:
                if grad_img[h][v] >= grad_img[h - 1][v - 1] and grad_img[h][v] >= grad_img[h + 1][v + 1]:
                    z[h][v] = grad_img[h][v]
        except IndexError as exc:
            # Handle boundary index errors
            pass

    return z

    @staticmethod
    def __angle__(a):
        """
        Convert gradient directions in radian to 4 possible direction in degree system

        :param a: Radian value of gradient direction numpy ndarray matrix
        :return: A int within {0, 45, 90, 135}
        """

        angle = np.rad2deg(a) % 180
        if (0 <= angle < 22.5) or (157.5 <= angle < 180):
            angle = 0
        elif 22.5 <= angle < 67.5:
            angle = 45
        elif 67.5 <= angle < 112.5:
            angle = 90
        elif 112.5 <= angle < 157.5:
            angle = 135
        return angle

```

```


In [21]: non_max_suppression = NonMaxSuppression()
         image_non_max = non_max_suppression(image_grad, image_theta)

```

```

# plotting
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20, 15))
ax[0].imshow(image_grad, cmap='gray')
ax[1].imshow(image_non_max, cmap='gray')
ax[0].set_title('grad image')
ax[1].set_title('non max suppressed')
plt.show()

```



1.1.4 1.D Thresholding

```

In [8]: class Thresholding:
    def __init__(self, high_threshold = 90, low_threshold = 30):
        self.high_threshold = high_threshold
        self.low_threshold = low_threshold
        self.weak = 29
        self.strong = 255
        self.flag = self.weak*9

    def _threshold_image(self, image):
        thresholded = np.empty(image.shape)
        thresholded[np.where(image>self.high_threshold)] = self.strong
        thresholded[np.where(((image>self.low_threshold) & (image<=self.high_threshold)))] = self.weak
        return thresholded

    def __call__(self, image):
        thresholded = self._threshold_image(image)
        for i in range(thresholded.shape[0]):
            for j in range(thresholded.shape[1]):

```



```

        if thresholded[i, j] == self.weak:
            if np.sum(thresholded[i-1:i+2, j-1:j+2]) > self.flag:
                thresholded[i, j] = self.strong
            else:
                thresholded[i, j] = 0
    return thresholded

```

```

In [32]: thresholding = Thresholding()
         thresholded = thresholding(image_non_max)


```

```

# plotting
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20, 15))
ax[0].imshow(image_non_max, cmap='gray')
ax[1].imshow(thresholded, cmap='gray')
ax[0].set_title('non max suppressed')
ax[1].set_title('thresholded')

plt.show()

```



```

In [26]: image = cv2.imread('images/cameraman.jpg', 0).astype(float)
         image_blurred = GaussianNoise()(image)
         image_grad, image_theta = GradientIntensity()(image_blurred)
         image_suppressed = NonMaxSuppression()(image_grad, image_theta)
         image_final = Thresholding()(image_suppressed)

```


```

# plotting
fig, ax = plt.subplots(nrows=1, ncols=2, figsize=(20, 15))
ax[0].imshow(image, cmap='gray')
ax[1].imshow(image_final, cmap='gray')

```

```
ax[0].set_title('input')
ax[1].set_title('canny edges')


plt.show()
```



1.2 2 Straight Line Detector

1. Assign each edge to a direction
2. Getting edgelets using connected components
3. Compute straightness and theta
4. Threshold
5. Test

```
In [52]: image = np.ones((10, 10), dtype=np.float) * 255
         image[:, 5:] = 0
         plt.imshow(image, cmap='gray')
         image_blurred = GaussianNoise()(image)
         image_grad, image_theta = GradientIntensity()(image_blurred)
         image_suppressed = NonMaxSuppression()(image_grad, image_theta)
         image_final = Thresholding()(image_suppressed)
         image_final[image_final <= 0] = 0
         image_final[image_final > 0] = 255
         image_final = image_final.astype(np.uint8)
```




1.2.1 2.A Assign each edge to a direction

```
In [53]: def assign_direction(grad_directions):
         """
         Assign each gradian direction to one of 4 discrete values

         :param grad_directions: theta values of gradients of canny edge detector
         :return: Directions using degrees only in 4 discrete values
         """
         angle = np.rad2deg(grad_directions) % 180
         angle[np.logical_or(np.logical_and(angle < 22.5, 0 <= angle),
                             np.logical_and(angle <= 180, 157.5 <= angle))] = 0
         angle[np.logical_and(angle < 67.5, 22.5 <= angle)] = 45
         angle[np.logical_and(angle < 112.5, 67.5 <= angle)] = 90
         angle[np.logical_and(angle < 157.5, 112.5 <= angle)] = 135
         return angle.astype(np.uint8)
         directions = assign_direction(image_theta)
         plt.imshow(directions, cmap='gray')
```

Out[53]: <matplotlib.image.AxesImage at 0x1c5e325ebe0>



1.2.2 2.B Getting edgelets using connected components

In [54]: `from scipy.ndimage import measurements`


`ANGLES = [0, 45, 90, 135]`

```
def get_connected_components(directions, angles):
    """
    Computes connected components of each direction in angles

    :param direction: Direction of gradients matrix
    :param angles: possible angles of given direction matrix
    """
    structure = np.array([[0,1,0], [1,1,1], [0,1,0]]) # + form connected components
    connected_components = {}
    for a in ANGLES:
        connected_components[str(a)] = measurements.label((directions == a)*1)[0]
    return connected_components

connected_components = get_connected_components(directions, ANGLES)
plt.imshow(connected_components['90'], cmap='gray')
```

Out [54]: `<matplotlib.image.AxesImage at 0x1c5e32b28d0>`



1.2.3 2.C Compute straightness and theta

```
In [55]: def raw_moment(data, iord, jord):
    nrows, ncols = data.shape
    y, x = np.mgrid[:nrows, :ncols]
    data = data * x**iord * y**jord
    return data.sum()

def inertial_axis(data):
    """
    Calculates the covariance matrix of image

    :param image: data
    :return: covariance matrix
    """
    data_sum = data.sum()
    m10 = raw_moment(data, 1, 0)
    m01 = raw_moment(data, 0, 1)
    x_bar = m10 / (data_sum+1e-10)
    y_bar = m01 / (data_sum+1e-10)
    u11 = (raw_moment(data, 1, 1) - x_bar * m01) / (data_sum+1e-10)
    u20 = (raw_moment(data, 2, 0) - x_bar * m10) / (data_sum+1e-10)
    u02 = (raw_moment(data, 0, 2) - y_bar * m01) / (data_sum+1e-10)
    cov = np.array([[u20, u11], [u11, u02]])
    return cov
```

```

def get_cov_matrices(connected_components, angles):
    """
    Computes covariance matrix of given connected components of direction of gradient.

    :param connected_components: connected_components dictionary of gradient direction
    :param angles: possible angles of given direction matrix
    """

    cov_matrices = {}
    for a in ANGLES:
        n_components = np.max(connected_components[str(a)])
        cov_comp = np.zeros((n_components, 2, 2))
        for i in range(n_components):
            cov_comp[i, :, :] = intertial_axis(connected_components[str(a)]==i)
        cov_matrices[str(a)] = cov_comp
    return cov_matrices

cov_matrices = get_cov_matrices(connected_components, ANGLES)
cov_matrices['90']

Out[55]: array([[ 2.50000000e-01,  1.12496679e-11],
                [ 1.12496679e-11,  8.25000000e+00]])

In [56]: def get_eigens(cov_matrices, angles):
    """
    Computes eigen vector and eigen values of covariance matrices for each direction

    :param cov_matrices: covariance matrices of connected components
    :param angles: possible angles of given direction matrix
    """

    eigvals = {}
    eigvecs = {}
    for a in ANGLES:
        n_components = cov_matrices[str(a)].shape[0]
        eigvals_comp = np.zeros((n_components, 2,))
        eigvecs_comp = np.zeros((n_components, 2, 2))
        for i in range(n_components):
            eigvals_comp[i, :], eigvecs_comp[i, :, :] = np.linalg.eigh(cov_matrices[str(a)]
            eigvals[str(a)], eigvecs[str(a)] = eigvals_comp, eigvecs_comp
    return eigvals, eigvecs

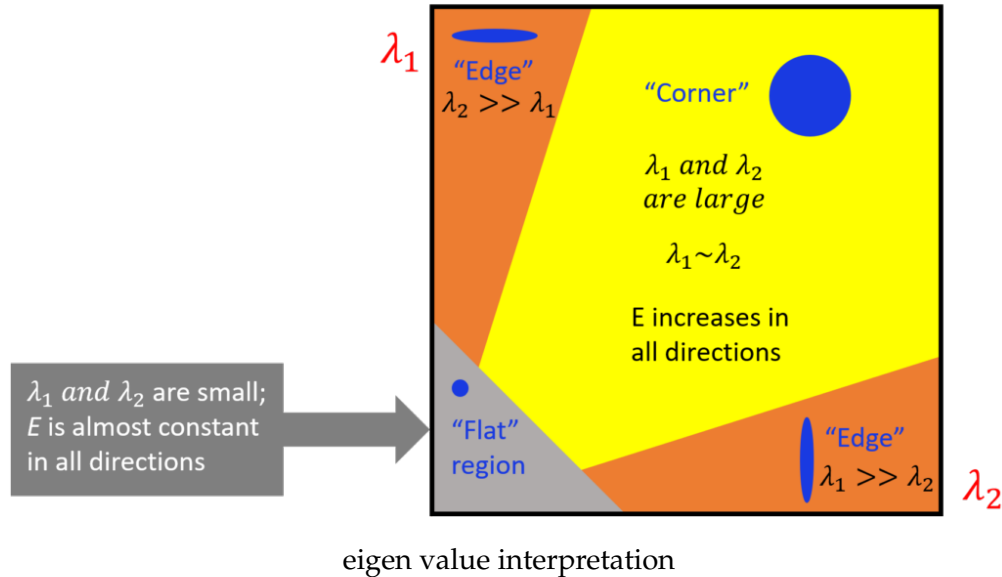
eigvals, eigvecs = get_eigens(cov_matrices, ANGLES)

In [57]: def get_theta_conf(eigvals, eigvecs, angles):
    """
    Computes theta(direction) and confidence degrees of lines using eigen matrices

```

Interpreting the eigenvalues

Classification of image points using eigenvalues of M :



```
:param eignvecs: Eigen vectors of covariance matrix of each direction
:param eignvals: Eigen values of covariance matrix of each direction
:param angles: possible angles of given direction matrix
"""
```

```
thetas = {}
confidences = {}
```

```
for a in ANGLES:
    n_components = eigvals[str(a)].shape[0]
    thetas_comp = np.zeros((n_components, 1))
    confidences_comp = np.zeros((n_components, 1))
    for i in range(n_components):
        thetas_comp[i,:] = np.rad2deg(np.arctan2(eigvecs[str(a)][i][1,1], eigvecs
            confidences_comp[i,:] = eigvals[str(a)][i][1] / (eigvals[str(a)][i][0]+1e
        thetas[str(a)] = thetas_comp
        confidences[str(a)] = confidences_comp
    return thetas, confidences
```

```
thetas, confidences = get_theta_conf(eigvals, eigvecs, ANGLES)
print(thetas['90'][0])
print(confidences['90'][0])
```

89.9999999999

32.999999987

1.2.4 2.D Threshold

```
In [58]: threshold = 5
```

```
def threshold_conf(confidences, threshold, angles):
    """
    Thresholds each confidence of each direction for all directions in angles

    :param confidences: confidence degrees of lines using eigen matrices
    :param threshold: threshold value for confidence score
    :param angle: possible angles of given direction matrix
    """

    strong_confs = {}
    for a in ANGLES:
        n_components = confidences[str(a)].shape[0]
        strong_conf_comp = np.zeros((n_components, 1))
        for i in range(n_components):
            strong_conf_comp[i,:] = 1 if confidences[str(a)][i] > threshold else 0
        strong_confs[str(a)] = strong_conf_comp
    return strong_confs


strong_confs = threshold_conf(confidences, threshold, ANGLES)
strong_confs
```

```
Out [58]: {'0': False, '45': False, '90': True, '135': False}
```

1.2.5 2.E Test

```
In [59]: image = cv2.imread('images/cameraman.jpg', 0).astype(float)
image_blurred = GaussianNoise()(image)
image_grad, image_theta = GradientIntensity()(image_blurred)
image_suppressed = NonMaxSuppression()(image_grad, image_theta)
image_final = Thresholding()(image_suppressed)
image_final[image_final < 1] = 0
image_final[image_final > 1] = 255
image_final = image_final.astype(np.uint8)
plt.imshow(image_final, cmap='gray')
```


```
Out [59]: <matplotlib.image.AxesImage at 0x1c5e330d390>
```

```
In [60]: from copy import deepcopy
         image_ = deepcopy(image_final)
```

```
In [121]: directions = assign_direction(image_theta)
          directions = np.multiply(directions, image_final==255) # mask out non-edge directions
          plt.imshow(directions, cmap='gray')
```

```
Out[121]: <matplotlib.image.AxesImage at 0x1c5e4afd240>
```




```

In [85]: ANGLES = [0, 45, 90, 135]
         connected_components = get_connected_components(directions, ANGLES)

         plt.figure(figsize=(20, 20))

         for i,a in enumerate(ANGLES):
             plt.subplot(1,4, int(i+1))
             plt.imshow(connected_components[str(a)], cmap='gray')
             plt.title('#comp'+str(np.max(connected_components[str(a)])))

```



```

In [123]: cov_matrices = get_cov_matrices(connected_components, ANGLES)
          eigvals, eigvecs = get_eigens(cov_matrices, ANGLES)

In [125]: thetas, confidences = get_theta_conf(eigvals, eigvecs, ANGLES)

         print('thetas:')
         for a in ANGLES:
             print(thetas[str(a)].shape)
         print('\n\nconfidences:')
         for a in ANGLES:
             print(confidences[str(a)].shape)

```

```

thetas:
(1, 1)
(163, 1)
(141, 1)
(137, 1)

```

```

confidences:
(1, 1)
(163, 1)
(141, 1)
(137, 1)

```

```

In [255]: threshold = np.mean([np.mean(confidences[str(a)]) for a in ANGLES])/10
          strong_confs = threshold_conf(confidences, threshold, ANGLES)

          for a in ANGLES:
              print('# of connected components with enough conf for {} degree: {}'.format(a, i))

# of connected components with enough conf for 0 degree: 0
# of connected components with enough conf for 45 degree: 8
# of connected components with enough conf for 90 degree: 61
# of connected components with enough conf for 135 degree: 4

In [256]: connected_components['90'].shape

Out[256]: (256, 256)


In [257]: strong_confs['90'].nonzero()[0]

Out[257]: array([ 1,  2,  3,  4,  5,  6,  7, 10, 11, 15, 19, 20, 21,
                  22, 25, 27, 28, 29, 31, 34, 36, 37, 38, 39, 43, 46,
                  47, 48, 50, 51, 52, 53, 55, 57, 59, 61, 68, 75, 76,
                  77, 78, 81, 82, 91, 92, 93, 94, 95, 96, 97, 100, 103,
                  108, 112, 116, 120, 123, 124, 125, 127, 135], dtype=int64)

In [303]: plt.imshow(connected_components['90']==46, cmap='gray')

Out[303]: <matplotlib.image.AxesImage at 0x1c5e96e4e10>

```



```
In [305]: plt.imshow(connected_components['45']==51, cmap='gray')
```

```
Out[305]: <matplotlib.image.AxesImage at 0x1c5e9796f28>
```

